Zero-Shot and Translation Experiments on XQuAD, MLQA and TyDiQA

Julen Etxaniz
UPV/EHU
jetxaniz007@ikasle.ehu.eus

Oihane Cantero
UPV/EHU
ocantero003@ikasle.ehu.eus

Abstract

1 Introduction

In this project we performed some zero-shot and translation experiments on Multilingual Question Answering. The objective is to compare the results of zero shot, translation test and translation test on different datasets, with different models. The datasets we used are XQuAD, MLQA and TyDiQA, and the models are monolingual or multilingual:

- Monolingual models:
  1. BERT (110M)
  2. BERT-large (340M)
  3. RoBERTa
  4. RoBERTa-large
- Multilingual models
  1. mBERT (110M)
  2. XLM-R
  3. XLM-R-large

Most of the models we used are already fine-tuned and available on Huggingface.

2 Related Work

3 Data

3.1 XQuAD

XQuAD is a multilingual Question Answering dataset (Artetxe et al., 2019). It is composed of 240 paraghaps and 1190 question-answer pair from SQuAD v1.1[1]. SQuAD is based on a set of Wikipedia articles. Professional translations into 11 languages were added in XQuAD (Spanish, German, Greek, Russian, Turkish, Arabic, Vietnamese, Thai, Chinese, Hindi and Romanian). As the dataset is based on SQuAD v1.1, there are no unanswerable questions in the data.

We also used XTREME (Hu et al., 2020) for automatically translated translate-train and translate-test data. The dataset can be found in HuggingFace [2].

3.2 MLQA

MLQA (Lewis et al., 2019) is another multilingual question answering evaluation benchmark. It has 5K extractive question-answering instances (12K in English) in seven languages (English, Arabic, German, Spanish, Hindi, Vietnamese and Simplified Chinese). It is also based on Wikipedia articles, and the questions has been translated by professional translators, while he answers are directly taken form the different languages of the given Wikipedia article, to get parallel sentences.

We also used XTREME (Hu et al., 2020) for automatically translated translate-train and translate-test data. The dataset can be found in HuggingFace [3].

3.3 TyDiQA

TyDi QA is a question answering dataset covering 11 typologically diverse languages with 204K question-answer pairs (Clark et al., 2020). The languages of TyDi QA are diverse with regard to their typology – the set of linguistic features that each language expresses – such that we expect models performing well on this set to generalize across a large number of the languages in the world. It contains language phenomena that would not be found in English-only corpora. To provide a realistic information-seeking task and avoid priming effects, questions are written by people who want to know the answer, but don't know the answer yet, (unlike SQuAD and its descendents) and the data

[1] https://huggingface.co/datasets/squad

[2] https://huggingface.co/datasets/juletxara/xquad_xtreme

[3] https://huggingface.co/datasets/mlqa

is collected directly in each language without the use of translation (unlike MLQA and XQuAD).

We also used XTREME (Hu et al., 2020) for automatically translated translate-train and translate-test data. The dataset can be found in HuggingFace [4].

4 Methods

All the code can be found on GitHub [5]

4.1 Zero-shot

4.2 Translate Train

4.3 Translate Test Monolingual

4.4 Translate Test Multilingual

4.5 Fine tuning

4.6 Data augmentation

5 Results

5.1 XQuAD

The results are obtained with XQuAD dataset are in Table 1. They are quite similar to those form the baseline and these are some conclusions we got.

We can see that zero-shot is better than translate-test for larger models and worse for smaller models. So we can deduce that larger models have more adaptability to unseen languages than smaller ones. Monolingual models get better results that multilingual ones translate-test, and as we might expect, larger models give better results than smaller ones.

Overall, the results from worst to better have been: Translate-train, Translate-test multilingual, monolingual, zero-shot, data augmentation and fine-tuning. The comparison of the results with fine tuning and data augmentation is not very pertinent because the fine tuning has been done with a part of the testing data. As we don't know which part has been used to fine tune the models, we couldn't remove them from the testing data.

The best languages have been English, Spanish, Romanian and Russian and the worst ones have been Chinese, Hindi, Thai and Turkish, with some very bad results, as for example, 25.2 F1 score and 16.8 EM for fine tuned mBERT in Thai. This could be because these four languages are are not in Latin script, and because Thai is not supported by mBERT.

5.2 MLQA

In the MLQA dataset also, we get higher results with the biggest models. The language that obtains the best score is English. It is not unexpected because the dataset has much more data in English than in the other languages.

5.3 MLQA Zero-Shot

mbert
  xlm-r
  xlm-r-large

5.4 TyDiQA

6 Conclusions

References

Mikel Artetxe, Sebastian Ruder, and Dani Yogatama. 2019. On the cross-lingual transferability of monolingual representations. *CoRR*, abs/1910.11856.

Jonathan H. Clark, Eunsol Choi, Michael Collins, Dan Garrette, Tom Kwiatkowski, Vitaly Nikolaev, and Jennimaria Palomaki. 2020. Tydi qa: A benchmark for information-seeking question answering in typologically diverse languages. *Transactions of the Association for Computational Linguistics*.

Junjie Hu, Sebastian Ruder, Aditya Siddhant, Graham Neubig, Orhan Firat, and Melvin Johnson. 2020. Xtreme: A massively multilingual multi-task benchmark for evaluating cross-lingual generalization.

Patrick Lewis, Barlas Oğuz, Ruty Rinott, Sebastian Riedel, and Holger Schwenk. 2019. Mlqa: Evaluating cross-lingual extractive question answering. *arXiv preprint arXiv:1910.07475*.

[4] https://huggingface.co/datasets/juletxara/tydiqa_xtreme

[5] https://github.com/juletx/XQuAD-MLQA

Model F1 / EM	en	ar	de	el	es	hi	ru	th	tr	vi	zh	ro	avg
Zero-shot													
mBERT	85.0 / 73.5	57.8 / 42.2	72.6 / 55.9	62.2 / 45.2	76.4 / 58.1	55.3 / 40.6	71.3 / 54.7	35.1 / 26.3	51.1 / 34.9	68.1 / 47.9	58.2 / 47.3	72.4 / 59.5	63.8 / 48.8
XLM-R	84.4 / 73.8	67.9 / 52.1	75.3 / 59.8	74.3 / 57.0	77.0 / 59.2	69.0 / 52.5	75.1 / 58.6	68.0 / 56.4	68.0 / 51.8	73.6 / 54.5	65.0 / 55.0	80.0 / 66.3	73.1 / 58.1
XLM-R Large	**86.5 / 75.9**	**75.0 / 58.0**	**79.9 / 63.8**	**79.1 / 61.3**	**81.0 / 62.7**	**76.0 / 60.8**	**80.3 / 63.1**	**72.8 / 61.7**	**74.1 / 58.3**	**79.0 / 59.3**	**66.8 / 58.0**	**83.5 / 70.2**	**77.8 / 62.8**
Translate-test monolingual													
BERT		69.4 / 55.0	75.7 / 62.7	75.0 / 60.6	77.2 / 62.6	69.7 / 53.7	74.9 / 60.5	60.5 / 46.5	59.9 / 41.8	72.2 / 58.3	69.9 / 56.0		70.4 / 55.8
BERT Large		73.6 / 59.1	80.4 / 66.4	80.2 / 66.8	81.9 / 68.7	**75.3 / 61.7**	80.1 / 67.0	**67.5 / 53.9**	**66.3 / 47.3**	**76.4 / 62.1**	74.0 / 59.5		75.6 / 61.2
RoBERTa		71.6 / 57.0	77.0 / 62.4	76.8 / 63.9	80.0 / 64.6	72.0 / 55.6	77.2 / 62.4	62.2 / 46.6	63.4 / 44.1	72.4 / 56.6	72.4 / 57.9		72.5 / 57.1
RoBERTa Large		**74.8 / 61.1**	**80.4 / 67.1**	**80.8 / 68.0**	**83.1 / 69.4**	75.1 / 61.0	**81.2 / 68.0**	65.3 / 51.0	66.0 / 46.9	76.4 / 62.0	**74.0 / 59.9**		**75.7 / 61.4**
Translate-test multilingual													
mBERT		70.4 / 55.8	76.7 / 63.3	76.0 / 61.9	78.7 / 65.1	70.6 / 55.8	76.6 / 63.1	60.0 / 45.9	61.6 / 42.7	70.6 / 55.6	70.1 / 56.6		71.2 / 56.6
XLM-R		70.4 / 56.5	79.0 / 65.8	77.8 / 65.0	79.3 / 66.4	72.4 / 57.6	77.4 / 63.6	60.3 / 45.4	63.4 / 44.3	73.0 / 58.4	71.1 / 57.4		72.4 / 58.0
XLM-R Large		**72.9 / 59.1**	**80.1 / 66.6**	**79.6 / 66.2**	**81.5 / 67.1**	**74.2 / 60.1**	**79.7 / 65.7**	**61.7 / 46.0**	**66.2 / 48.2**	**75.1 / 61.5**	**73.6 / 58.8**		**74.5 / 59.9**
Translate-train													
XLM-R-es	**80.4** / 66.1	**67.0** / 47.9	74.2 / 56.4	**73.5**/ 52.4	**76.3** / 56.6	**66.9**/ 48.2	72.4 / 54.2	**68.7 / 58.5**	**66.2** / 46.5	73.2 / 52.0	63.4 / 50.3	**76.0** / 59.2	**71.5** / 54.0
XLM-R-de	79.8 / **67.1**	65.9 / **48.2**	**74.3 / 58.8**	72.3 / **54.4**	75.9 / **57.9**	66.4 / **50.6**	**73.1 / 56.4**	65.4 / 56.8	65.8 / **50.8**	72.7 / **53.2**	**64.7 / 55.0**	75.3 / **61.1**	71.0 / **55.9**
Fine-tuning XQuAD													
mBERT	97.3 / 95.3	90.0 / 84.3	94.2 / 90.0	92.2 / 87.0	96.2 / 92.4	88.2 / 77.5	94.4 / 90.1	25.2 / 16.8	89.9 / 84.4	93.4 / 87.6	87.5 / 84.4	95.5 / 91.3	87.0 / 81.8
XLM-R	98.5 / 97.5	92.5 / 88.2	95.1 / 91.8	96.0 / 91.8	97.8 / 93.6	92.6 / 88.6	95.2 / 90.8	94.0 / 92.4	92.0 / 87.3	95.5 / 91.3	94.0 / 92.9	97.7 / 94.8	95.1 / 91.8
XLM-R Large	**99.7 / 99.2**	**97.0 / 94.2**	**98.1 / 95.6**	**97.8 / 94.4**	**98.5 / 95.8**	**96.5 / 93.6**	**98.1 / 96.0**	**96.1 / 95.1**	**95.9 / 92.3**	**97.6 / 94.0**	**96.3 / 95.7**	**98.9 / 97.1**	**97.5 / 95.2**
Data-augmentation XQuAD													
mBERT	99.7 / 99.2	97.1 / 94.4	98.9 / 97.9	97.0 / 94.6	99.6 / 98.9	97.7 / 95.1	98.5 / 97.3	87.3 / 84.9	98.8 / 97.4	98.9 / 97.5	97.5 / 96.8	90.6 / 81.6	96.8 / 94.6

Table 1: XQuAD results (F1/EM) for each language.

Model F1 / EM	en	es	de	ar	hi	vi	zh	avg
Zero-shot								
mBERT	80.3 / 67.0	64.9 / 43.6	59.4 / 43.8	44.9 / 28.0	46.2 / 30.0	58.8 / 39.6	37.4 / 36.8	56.0 / 41.3
XLM-R	80.8 / 68.0	66.5 / 46.1	62.2 / 46.7	54.6 / 36.0	61.4 / 44.2	67.2 / 46.3	40.0 / 39.3	61.8 / 46.7
XLM-R Large	**84.0 / 71.2**	**72.1 / 50.2**	**68.5 / 52.4**	**62.0 / 42.1**	**69.8 / 51.3**	**73.1 / 51.8**	**45.7 / 45.1**	**67.9 / 52.0**
Translate-test monolingual								
BERT		65.0 / 43.2	54.4 / 35.7	51.0 / 27.7	52.8 / 32.0	53.6 / 32.1	47.8 / 26.6	54.1 / 32.9
BERT Large		67.2 / 45.2	56.7 / 37.2	52.7 / 28.9	55.2 / 33.8	**56.7** / 34.7	50.1 / **27.8**	56.4 / 34.6
RoBERTa		66.0 / 43.4	54.1 / 34.1	51.4 / 27.6	52.3 / 31.0	54.0 / 32.4	47.6 / 25.2	54.3 / 32.3
RoBERTa Large		**68.0 /45.9**	**57.4 / 38.0**	**53.7 / 29.4**	**55.7 / 33.9**	56.3 / **34.9**	**50.6** / 27.7	**56.9 / 35.0**
Translate-test multilingual								
mBERT		64.3 / 43.0	53.6 / 34.8	49.5 / 27.0	51.9 / 31.2	53.4 / 32.0	45.9 / 24.5	53.1 / 32.1
XLM-R		64.8 / 43.0	53.6 / 34.9	50.4 / 27.7	52.8 / 32.0	54.2 / 33.4	47.7 / 26.1	53.9 / 32.9
XLM-R Large		**68.6 / 46.5**	**56.6 / 37.4**	**53.1 / 29.2**	**55.6 / 34.5**	**56.6/ 34.5**	**50.0 / 27.6**	**56.7 / 35.0**
Translate-train								
XLM-R-es	77.2 / 61.5	**68.0** / 44.8	61.4 / 44.9	**54.1** / 34.1	**60.2** / 40.7	**66.2** / 45.0	36.2 / 35.4	**60.5** / 43.8
XLM-R-de	**77.3 / 63.6**	65.6 / **45.0**	**62.4 / 46.7**	53.6 / **35.6**	60.1 / **43.8**	65.0 / **45.2**	**38.1 / 37.4**	60.3 / **45.3**

Table 2: MLQA results (F1/EM) for each language.

c/q	en	es	de	ar	hi	vi	zh	avg
en	80.3 / 67.0	67.4 / 52.8	66.4 / 52.5	44.1 / 31.1	39.3 / 26.3	53.7 / 39.1	55.8 / 41.4	58.1 / 44.3
es	66.9 / 46.4	64.9 / 43.6	60.6 / 40.2	43.1 / 26.0	36.2 / 20.1	48.5 / 31.4	49.9 / 30.6	52.9 / 34.0
de	62.4 / 46.7	56.4 / 41.0	59.4 / 43.8	36.8 / 23.6	34.0 / 21.5	43.6 / 29.6	46.5 / 30.7	48.4 / 33.8
ar	51.1 / 33.7	45.4 / 28.7	46.3 / 30.5	44.9 / 28.0	30.8 / 17.3	35.9 / 20.1	36.8 / 21.3	41.6 / 25.7
hi	52.9 / 37.1	43.7 / 29.1	47.6 / 33.8	34.5 / 21.4	46.2 / 30.0	38.0 / 25.0	39.2 / 25.2	43.2 / 28.8
vi	64.5 / 44.8	53.9 / 37.5	53.7 / 36.6	32.5 / 19.3	35.1 / 19.7	58.8 / 39.6	50.3 / 32.3	49.8 / 32.8
zh	38.3 / 37.7	29.0 / 28.3	30.0 / 28.9	21.0 / 20.6	16.6 / 16.2	25.1 / 24.4	37.4 / 36.8	28.2 / 27.6
avg	59.5 / 44.8	51.5 / 37.3	52.0 / 38.0	36.7 / 24.3	34.0 / 21.6	43.4 / 29.9	45.1 / 31.2	46.0 / 32.4

Table 3: MLQA results (F1/EM) for each language in zero-shot with mBERT. Columns show question language, rows show context language.

c/q	en	es	de	ar	hi	vi	zh	avg
en	80.8 / 68.0	57.8 / 43.9	60.8 / 47.1	33.5 / 21.3	45.0 / 32.0	39.8 / 27.5	37.9 / 25.3	50.8 / 37.9
es	66.0 / 45.1	66.5 / 46.1	50.5 / 32.6	25.2 / 12.3	31.8 / 17.1	29.1 / 14.9	28.2 / 14.3	42.5 / 26.1
de	60.0 / 44.3	44.0 / 29.7	62.2 / 46.7	22.2 / 12.1	29.4 / 17.6	28.7 / 16.2	29.1 / 17.2	39.4 / 26.3
ar	51.5 / 33.8	27.0 / 13.5	34.2 / 19.8	54.6 / 36.0	15.6 / 5.8	15.0 / 5.7	14.1 / 5.1	30.3 / 17.1
hi	60.6 / 43.4	37.4 / 23.0	42.8 / 27.8	19.5 / 8.0	61.4 / 44.2	24.3 / 11.9	26.1 / 13.6	38.9 / 24.6
vi	63.6 / 44.6	32.6 / 19.1	41.9 / 25.7	17.8 / 6.6	29.2 / 15.0	67.2 / 46.3	27.4 / 13.8	40.0 / 24.4
zh	34.9 / 34.3	11.3 / 10.7	14.0 / 13.3	3.9 / 3.7	10.8 / 10.4	8.1 / 7.7	40.0 / 39.3	17.6 / 17.1
avg	59.6 / 44.8	39.5 / 26.6	43.8 / 30.4	25.2 / 14.3	31.9 / 20.3	30.3 / 18.6	29.0 / 18.4	37.0 / 24.8

Table 4: MLQA results (F1/EM) for each language in zero-shot with XLM-R. Columns show question language, rows show context language.

c/q	en	es	de	ar	hi	vi	zh	avg
en	84.0 / 71.2	77.2 / 64.2	77.7 / 65.1	32.4 / 22.1	43.6 / 30.7	61.6 / 48.5	33.8 / 21.1	58.6 / 46.1
es	72.1 / 50.3	72.1 / 50.2	70.0 / 48.6	33.0 / 17.8	42.1 / 26.1	54.8 / 35.6	36.5 / 20.5	54.4 / 35.6
de	67.7 / 51.7	65.3 / 49.6	68.5 / 52.4	31.2 / 19.5	36.2 / 21.9	50.7 / 34.3	32.3 / 18.9	50.3 / 35.5
ar	61.7 / 42.2	56.7 / 38.4	59.7 / 41.8	62.0 / 42.1	43.9 / 27.4	48.6 / 30.7	38.7 / 21.6	53.0 / 34.9
hi	70.5 / 52.6	63.2 / 45.9	65.1 / 49.9	45.5 / 29.1	69.8 / 51.3	54.4 / 37.6	44.6 / 28.5	59.0 / 42.1
vi	72.1 / 50.9	64.7 / 45.5	67.7 / 48.2	35.8 / 20.9	42.0 / 25.3	73.1 / 51.8	39.2 / 21.7	56.4 / 37.8
zh	44.2 / 43.6	36.7 / 36.1	41.1 / 40.2	25.9 / 25.4	30.0 / 29.5	35.0 / 34.6	45.7 / 45.1	36.9 / 36.4
avg	67.5 / 51.8	62.3 / 47.1	64.3 / 49.5	38.0 / 25.3	43.9 / 30.3	54.0 / 39.0	38.7 / 25.3	52.7 / 38.3

Table 5: MLQA results (F1/EM) for each language in zero-shot with XLM-R-Large. Columns show question language, rows show context language.